

```
In [1]: import Pkg
Pkg.activate(@__DIR__)
Pkg.instantiate()
using LinearAlgebra, Plots
import ForwardDiff as FD
import MeshCat as mc
using Test
```

**Activating** project at `c:\Users\zsqu4\Desktop\OCRL\Optimal-Control-and-Reinforcement-Learning\HW1\_S25`

## Julia Warmup

Just like Python, Julia lets you do the following:

```
In [2]: let
    x = [1,2,3]
    @show x
    y = x # NEVER DO THIS, EDITING ONE WILL NOW EDIT BOTH

    y[3] = 100 # this will now modify both y and x
    x[1] = 300 # this will now modify both y and x

    @show y
    @show x
end

x = [1, 2, 3]
y = [300, 2, 100]
x = [300, 2, 100]
3-element Vector{Int64}:
 300
   2
 100
```

```
In [3]: # to avoid this, here are two alternatives
let
    x = [1,2,3]
    @show x

    y1 = 1*x          # this is fine
    y2 = deepcopy(x) # this is also fine

    x[2] = 200 # only edits x
    y1[1] = 400 # only edits y1
    y2[3] = 100 # only edits y2

    @show x
    @show y1
    @show y2
end
```

```
x = [1, 2, 3]
x = [1, 200, 3]
y1 = [400, 2, 3]
y2 = [1, 2, 100]
3-element Vector{Int64}:
 1
 2
100
```

## Optional function arguments

We can have optional keyword arguments for functions in Julia, like the following:

```
In [4]: ## optional arguments in functions

# we can have functions with optional arguments after a ; that have default values
let
    function f1(a, b; c=4, d=5)
        @show a,b,c,d
    end

    f1(1,2)           # this means c and d will take on default value
    f1(1,2;c = 100,d = 2) # specify c and d
    f1(1,2;d = -30)     # or we can only specify one of them
end

(a, b, c, d) = (1, 2, 4, 5)
(a, b, c, d) = (1, 2, 100, 2)
(a, b, c, d) = (1, 2, 4, -30)
(1, 2, 4, -30)
```

## Q1: Integration (25 pts)

In this question we are going to integrate the equations of motion for a double pendulum using multiple explicit and implicit integrators. We will write a generic simulation function for each of the two categories (explicit and implicit), and compare 6 different integrators.

The continuous time dynamics of the cartpole are written as a function:

$$\dot{x} = f(x)$$

In the code you will see `xdot = dynamics(params, x)`.

### Part A (10 pts): Explicit Integration

Here we are going to implement the following explicit integrators:

- Forward Euler (explicit)
- Midpoint (explicit)
- RK4 (explicit)

```
In [5]: # these two functions are given, no TODO's here
function double_pendulum_dynamics(params::NamedTuple, x::Vector)
    # continuous time dynamics for a double pendulum given state x,
    # also known as the "equations of motion".
    # returns the time derivative of the state, ḡ (dx/dt)

    # the state is the following:
    θ₁, θ̇₁, θ₂, θ̇₂ = x

    # system parameters
    m₁, m₂, L₁, L₂, g = params.m₁, params.m₂, params.L₁, params.L₂, params.g

    # dynamics
    c = cos(θ₁ - θ₂)
    s = sin(θ₁ - θ₂)

    ḡ = [
        θ̇₁;
        (m₂ * g * sin(θ₂) * c - m₂ * s * (L₁ * c * θ₁^2 + L₂ * θ₂^2) - (m₁ + m₂) * g * sin(θ₁)) / (L₁
        θ̇₂;
        ((m₁ + m₂) * (L₁ * θ₁^2 * s - g * sin(θ₂) + g * sin(θ₁) * c) + m₂ * L₂ * θ₂^2 * s * c) / (L₂ * (m₁ + m₂))
    ]

    return ḡ
end
function double_pendulum_energy(params::NamedTuple, x::Vector)::Real
    # calculate the total energy (kinetic + potential) of a double pendulum given a state vector x

    # the state is the following:
    θ₁, θ̇₁, θ₂, θ̇₂ = x

    # system parameters
    m₁, m₂, L₁, L₂, g = params.m₁, params.m₂, params.L₁, params.L₂, params.g

    # cartesian positions/velocities of the masses
    r₁ = [L₁ * sin(θ₁), 0, -params.L₁ * cos(θ₁) + 2]
    r₂ = r₁ + [params.L₂ * sin(θ₂), 0, -params.L₂ * cos(θ₂)]
    v₁ = [L₁ * θ̇₁ * cos(θ₁), 0, L₁ * θ̇₁ * sin(θ₁)]
    v₂ = v₁ + [L₂ * θ̇₂ * cos(θ₂), 0, L₂ * θ̇₂ * sin(θ₂)]

    # energy calculation
    kinetic = 0.5 * (m₁ * v₁' * v₁ + m₂ * v₂' * v₂)
    potential = m₁ * g * r₁[3] + m₂ * g * r₂[3]
    return kinetic + potential
end
```

double\_pendulum\_energy (generic function with 1 method)

Now we are going to simulate this double pendulum by integrating the equations of motion with the simplest explicit integrator, the Forward Euler method:

$$x_{k+1} = x_k + \Delta t \cdot f(x_k) \quad \text{Forward Euler (explicit)}$$

```
In [6]: """
    x_{k+1} = forward_euler(params, dynamics, x_k, dt)

Given `ẋ = dynamics(params, x)`, take in the current state `x` and integrate it for
using Forward Euler method.
"""

function forward_euler(params::NamedTuple, dynamics::Function, x::Vector, dt::Real)

    for i in 1:length(x) - 1
        ẋ = dynamics(params, x[i])
        x[i+1] = x[i] + dt * ẋ
    end
    return x
end

forward_euler
```

```
In [7]: include(joinpath(@__DIR__, "animation.jl"))
using Test

@testset "forward euler" begin

    # parameters for the simulation
    params = (
        m1 = 1.0,
        m2 = 1.0,
        L1 = 1.0,
        L2 = 1.0,
        g = 9.8
    )

    # initial condition
    x0 = [pi/1.6; 0; pi/1.8; 0]

    # time step size (s)
    dt = 0.01
    tf = 30.0
    t_vec = 0:dt:tf
    N = length(t_vec)

    # store the trajectory in a vector of vectors
    X = [zeros(4) for i = 1:N]
    X[1] = 1*x0

    # TODO: simulate the double pendulum with `forward_euler`
    # X[k] = `x_k`, so X[k+1] = forward_euler(params, double_pendulum_dynamics, X[k])

    X = forward_euler(params, double_pendulum_dynamics, X, dt)

    # calculate energy
    E = [double_pendulum_energy(params, x) for x in X]

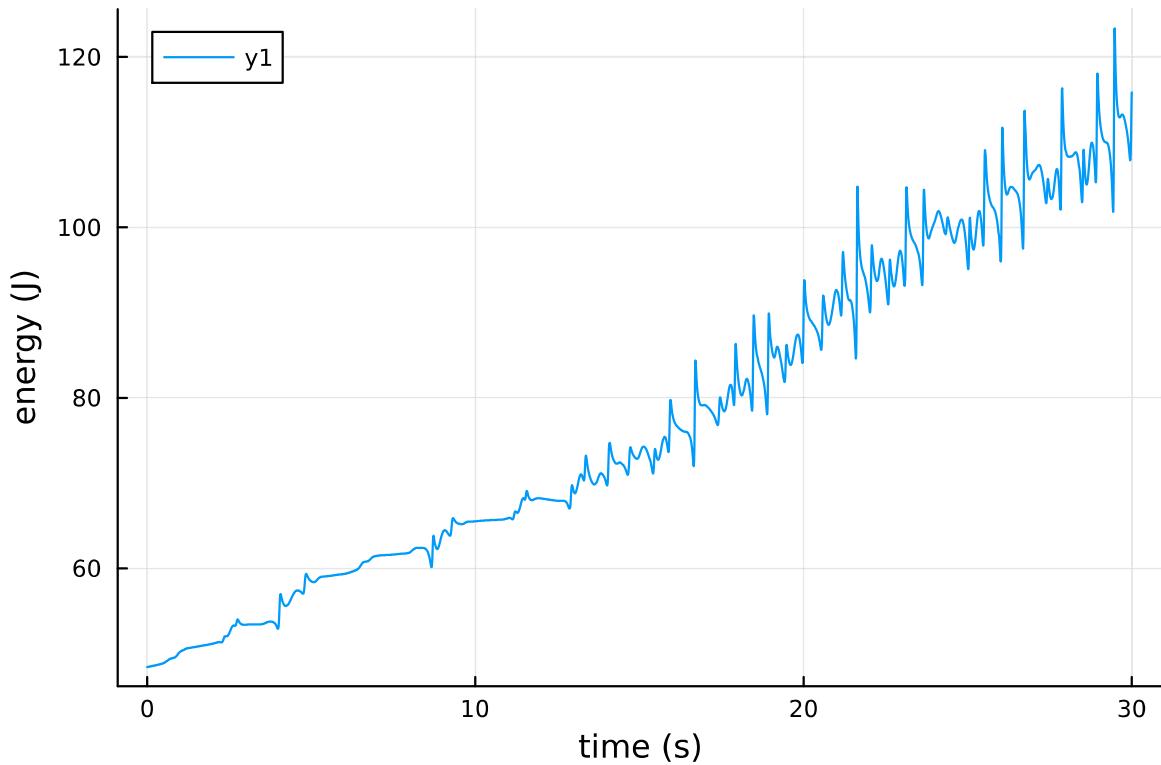
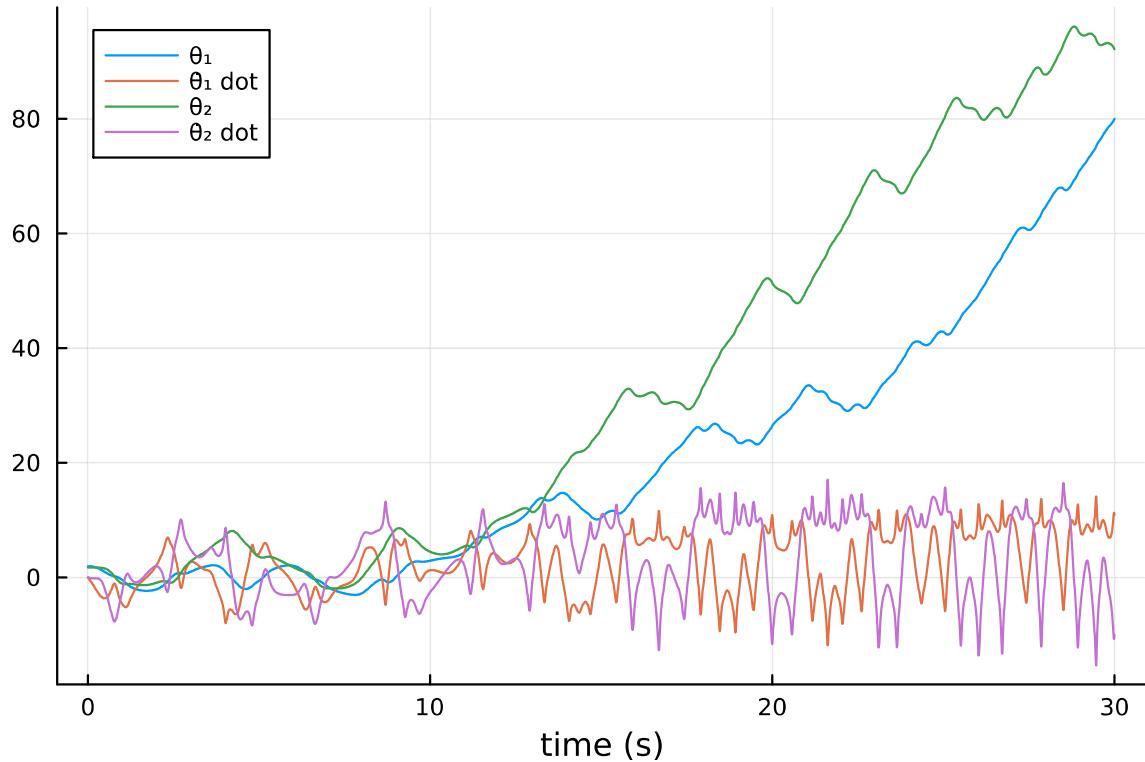
    @test norm(X[end]) > 1e-10    # make sure all X's were updated
    @test 2 < (E[end]/E[1]) < 3    # energy should be increasing

    # plot state history, energy history, and animate it
```

```

display(plot(t_vec, hcat(X...)', xlabel = "time (s)", label = ["θ1" "θ1 dot" "θ2" "θ2 dot"]))
display(plot(t_vec, E, xlabel = "time (s)", ylabel = "energy (J)"))
meshcat_animate(params,X,dt,N)
end

```



[ Info: Listening on: 127.0.0.1:8701, thread id: 1  
 [ @ HTTP.Servers C:\Users\zsqu4\.julia\packages\HTTP\4AUP1\src\Servers.jl:382

```
Test Summary: | Pass  Total  Time
forward euler |    2      2  6.8s
```

```
| Info: MeshCat server started. You can open the visualizer by visiting the following URL in your browser:
| http://127.0.0.1:8701
└ @ MeshCat C:\Users\zsqu4\.julia\packages\MeshCat\9QrxD\src\visualizer.jl:43
Test.DefaultTestSet("forward euler", Any[], 2, false, false, true, 1.738638680649e9,
1.738638687487e9, false, "c:\Users\zsqu4\Desktop\OCRL\Optimal-Control-and-Reinforcement-Learning\HW1_S25\jl_notebook_cell_df34fa98e69747e1a8f8a730347b8e2f_X14sZm1sZQ==.jl")
```

Now let's implement the next two integrators:

### Midpoint:

$$x_m = x_k + \frac{\Delta t}{2} \cdot f(x_k) \quad (1)$$

$$x_{k+1} = x_k + \Delta t \cdot f(x_m) \quad (2)$$

### RK4:

$$k_1 = \Delta t \cdot f(x_k) \quad (3)$$

$$k_2 = \Delta t \cdot f(x_k + k_1/2) \quad (4)$$

$$k_3 = \Delta t \cdot f(x_k + k_2/2) \quad (5)$$

$$k_4 = \Delta t \cdot f(x_k + k_3) \quad (6)$$

$$x_{k+1} = x_k + (1/6) \cdot (k_1 + 2k_2 + 2k_3 + k_4) \quad (7)$$

```
In [8]: function midpoint(params::NamedTuple, dynamics::Function, x::Vector, dt::Real)::Vector
    # TODO: implement explicit midpoint

    for i in 1:length(x)-1
        ḡ = dynamics(params, x[i])
        xm = x[i] + dt/2*ḡ
        x[i+1] = x[i] + dt*dynamics(params, xm)
    end
    # error("midpoint not implemented")
    # return 0*x
    return x
end
function rk4(params::NamedTuple, dynamics::Function, x::Vector, dt::Real)::Vector
    # TODO: implement RK4
    for i in 1:length(x)-1
        k1 = dt*dynamics(params, x[i])
        k2 = dt*dynamics(params, x[i] + k1/2)
        k3 = dt*dynamics(params, x[i] + k2/2)
        k4 = dt*dynamics(params, x[i] + k3)
        x[i+1] = x[i] + (k1 + 2*k2 + 2*k3 + k4)/6
    end
    return x
end
```

rk4 (generic function with 1 method)

```
In [9]: function simulate_explicit(params::NamedTuple,dynamics::Function,integrator::Function)
    # TODO: update this function to simulate dynamics forward
    # with the given explicit integrator

    # take in
    t_vec = 0:dt:tf
    N = length(t_vec)
    X = [zeros(length(x0)) for i = 1:N]
    X[1] = x0

    # TODO: simulate X forward
    X = integrator(params, dynamics, X, dt)
    # return state history X and energy E
    E = [double_pendulum_energy(params,x) for x in X]
    return X, E
end
```

simulate\_explicit (generic function with 1 method)

```
In [10]: # initial condition
const x0 = [pi/1.6; 0; pi/1.8; 0]

const params = (
    m1 = 1.0,
    m2 = 1.0,
    L1 = 1.0,
    L2 = 1.0,
    g = 9.8
)
```

(m1 = 1.0, m2 = 1.0, L1 = 1.0, L2 = 1.0, g = 9.8)

```
In [11]: # dt = 0.01
# tf = 30.0
# t_vec = 0:dt:tf
# N = Length(t_vec)
# x0 = [pi/1.6; 0; pi/1.8; 0]

# X = [zeros(4) for i = 1:N]
# X[1] = 1*x0

# X, E = simulate_explicit(params, double_pendulum_dynamics, midpoint, x0, dt, tf)

# display(plot(t_vec, hcat(X...)', xlabel = "time (s)", label = ["θ₁" "θ₁ dot" "θ₂"])
# display(plot(t_vec, E, xlabel = "time (s)", ylabel = "energy (J)"))
# meshcat_animate(params,X,dt,N)
```

## Part B (10 pts): Implicit Integrators

Explicit integrators work by calling a function with  $x_k$  and  $\Delta t$  as arguments, and returning  $x_{k+1}$  like this:

$$x_{k+1} = f_{\text{explicit}}(x_k, \Delta t)$$

Implicit integrators on the other hand have the following relationship between the state at  $x_k$  and  $x_{k+1}$ :

$$f_{\text{implicit}}(x_k, x_{k+1}, \Delta t) = 0$$

This means that if we want to get  $x_{k+1}$  from  $x_k$ , we have to solve for a  $x_{k+1}$  that satisfies the above equation. This is a rootfinding problem in  $x_{k+1}$  (our unknown), so we just have to use Newton's method.

Here are the three implicit integrators we are looking at, the first being Backward Euler (1st order):

$$f(x_k, x_{k+1}, \Delta t) = x_k + \Delta t \cdot \dot{x}_{k+1} - x_{k+1} = 0 \quad \text{Backward Euler}$$

Implicit Midpoint (2nd order)

$$x_{k+1/2} = \frac{1}{2}(x_k + x_{k+1}) \quad (8)$$

$$f(x_k, x_{k+1}, \Delta t) = x_k + \Delta t \cdot \dot{x}_{k+1/2} - x_{k+1} = 0 \quad \text{Implicit Midpoint} \quad (9)$$

Hermite Simpson (3rd order)

$$x_{k+1/2} = \frac{1}{2}(x_k + x_{k+1}) + \frac{\Delta t}{8}(\dot{x}_k - \dot{x}_{k+1})$$

$$f(x_k, x_{k+1}, \Delta t) = x_k + \frac{\Delta t}{6} \cdot (\dot{x}_k + 4\dot{x}_{k+1/2} + \dot{x}_{k+1}) - x_{k+1} = 0 \quad \text{Hermite-Simpson}$$

When you implement these integrators, you will update the functions such that they take in a dynamics function,  $x_k$  and  $x_{k+1}$ , and return the residuals described above. We are NOT solving these yet, we are simply returning the residuals for each implicit integrator that we want to be 0.

```
In [12]: # since these are explicit integrators, these function will return the residuals de
# NOTE: we are NOT solving anything here, simply return the residuals
function backward_euler(params::NamedTuple, dynamics::Function, x1::Vector, x2::Vec
    # error("backward euler not implemented")
    return x2 - x1 - dt*dynamics(params, x2)

end
function implicit_midpoint(params::NamedTuple, dynamics::Function, x1::Vector, x2::V
    # error("implicit midpoint not implemented")
    return x2 - x1 - dt*dynamics(params, (x1+x2)/2)
end
function hermite_simpson(params::NamedTuple, dynamics::Function, x1::Vector, x2::Ve
    # error("hermite simpson not implemented")
    x_05 = 0.5*(x1 + x2) + dt/8*(dynamics(params, x1) - dynamics(params, x2))
    return x2 - x1 - dt/6*(dynamics(params, x1) + 4*dynamics(params, x_05) + dynami
end

hermite_simpson (generic function with 1 method)
```

```
In [13]: # TODO
# this function takes in a dynamics function, implicit integrator function, and x1
# and uses Newton's method to solve for an x2 that satisfies the implicit integrator
# that we wrote about in the functions above
function implicit_integrator_solve(params::NamedTuple, dynamics::Function, implicit)

    # initialize guess
    x2 = 1*x1

    # TODO: use Newton's method to solve for x2 such that residual for the integrator
    # DO NOT USE A WHILE LOOP
    for i = 1:max_iters
        residual = x -> implicit_integrator(params, dynamics, x1, x, dt)
        R = residual(x2)
        J = FD.jacobian(residual,x2)
        x2_new = x2 - inv(J)*R
        # TODO: return x2 when the norm of the residual is below tol
        if norm(R) < tol
            return x2
        end
        x2 = x2_new
    end
    # error("implicit integrator solve failed")
end

implicit_integrator_solve (generic function with 1 method)
```

```
In [14]: @testset "implicit integrator check" begin

    dt = 1e-1
    x1 = [.1,.2,.3,.4]

    for integrator in [backward_euler, implicit_midpoint, hermite_simpson]
        println("-----testing $integrator -----")
        x2 = implicit_integrator_solve(params, double_pendulum_dynamics, integrator)
        @test norm(integrator(params, double_pendulum_dynamics, x1, x2, dt)) < 1e-1
    end
end

-----testing backward_euler -----
-----testing implicit_midpoint -----
-----testing hermite_simpson -----
Test Summary:           | Pass  Total  Time
implicit integrator check |    3      3  2.1s
Test.DefaultTestSet("implicit integrator check", Any[], 3, false, false, true, 1.738
638687932e9, 1.738638690036e9, false, "c:\\\\Users\\\\zsqu4\\\\Desktop\\\\OCRL\\\\Optimal-Cont
rol-and-Reinforcement-Learning\\\\HW1_S25\\\\jl_notebook_cell_df34fa98e69747e1a8f8a73034
7b8e2f_X26sZmlsZQ==.jl")
```

```
In [15]: function simulate_implicit(params::NamedTuple,dynamics::Function,implicit_integrator)
    t_vec = 0:dt:tf
    N = length(t_vec)
    X = [zeros(length(x0)) for i = 1:N]
    X[1] = x0

    # TODO: do a forward simulation with the selected implicit integrator
```

```
# hint: use your `implicit_integrator_solve` function
for i in 1:N-1
    X[i+1] = implicit_integrator_solve(params, dynamics, implicit_integrator, X)
end
# X = implicit_integrator_solve(params, dynamics, implicit_integrator, X[1], dt)

E = [double_pendulum_energy(params,x) for x in X]
@assert length(X)==N
@assert length(E)==N
return X, E
end

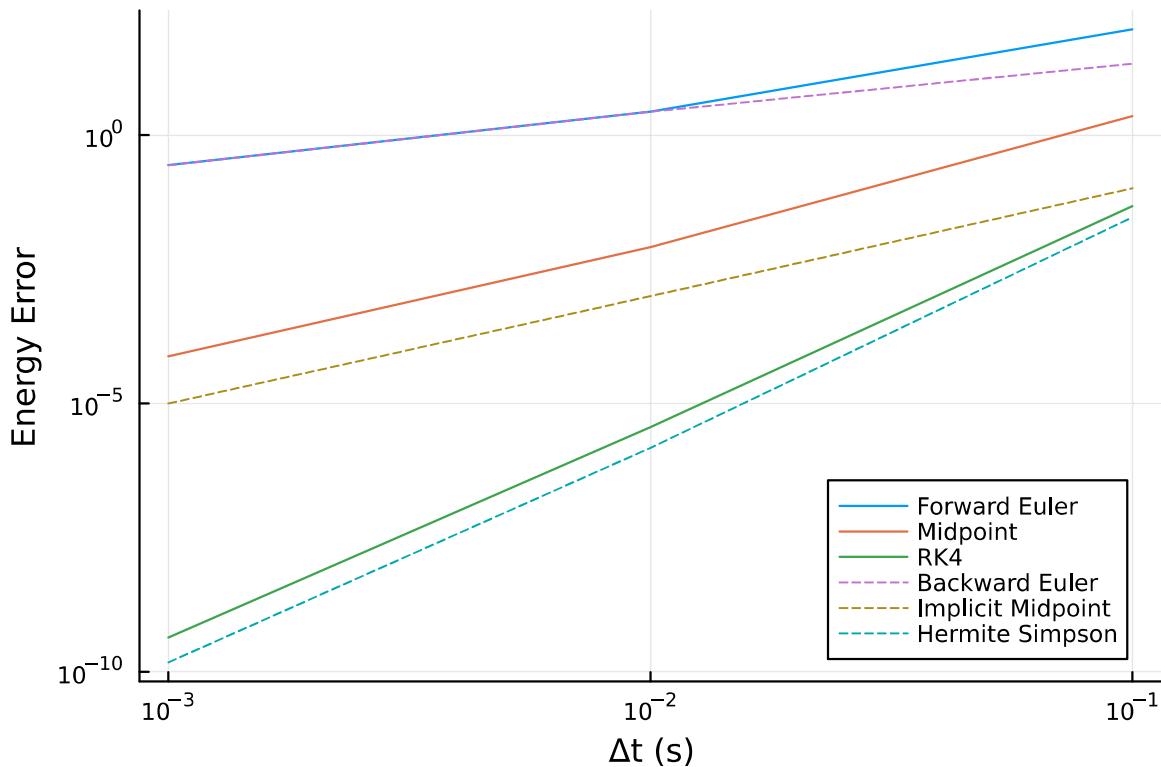
simulate_implicit (generic function with 1 method)
```

```
In [16]: function max_err_E(E)
    E0 = E[1]
    err = abs.(E .- E0)
    return maximum(err)
end
function get_explicit_energy_error(integrator::Function, dts::Vector)
    [max_err_E(simulate_explicit(params, double_pendulum_dynamics, integrator, x0, dt, t))
end
function get_implicit_energy_error(integrator::Function, dts::Vector)
    [max_err_E(simulate_implicit(params, double_pendulum_dynamics, integrator, x0, dt, t))
end

const tf = 2.0
let
    # here we compare everything
    dts = [1e-3, 1e-2, 1e-1]
    explicit_integrators = [forward_euler, midpoint, rk4]
    implicit_integrators = [backward_euler, implicit_midpoint, hermite_simpson]

    explicit_data = [get_explicit_energy_error(integrator, dts) for integrator in explicit_integrators]
    implicit_data = [get_implicit_energy_error(integrator, dts) for integrator in implicit_integrators]

    plot(dts, hcat(explicit_data...), label = ["Forward Euler" "Midpoint" "RK4"], xlabel = "dt")
    plot!(dts, hcat(implicit_data...), ls = :dash, label = ["Backward Euler" "Implicit Midpoint"])
    plot!(legend=:bottomright)
end
```



What we can see above is the maximum energy error for each of the integration methods. In general, the implicit methods of the same order are slightly better than the explicit ones.

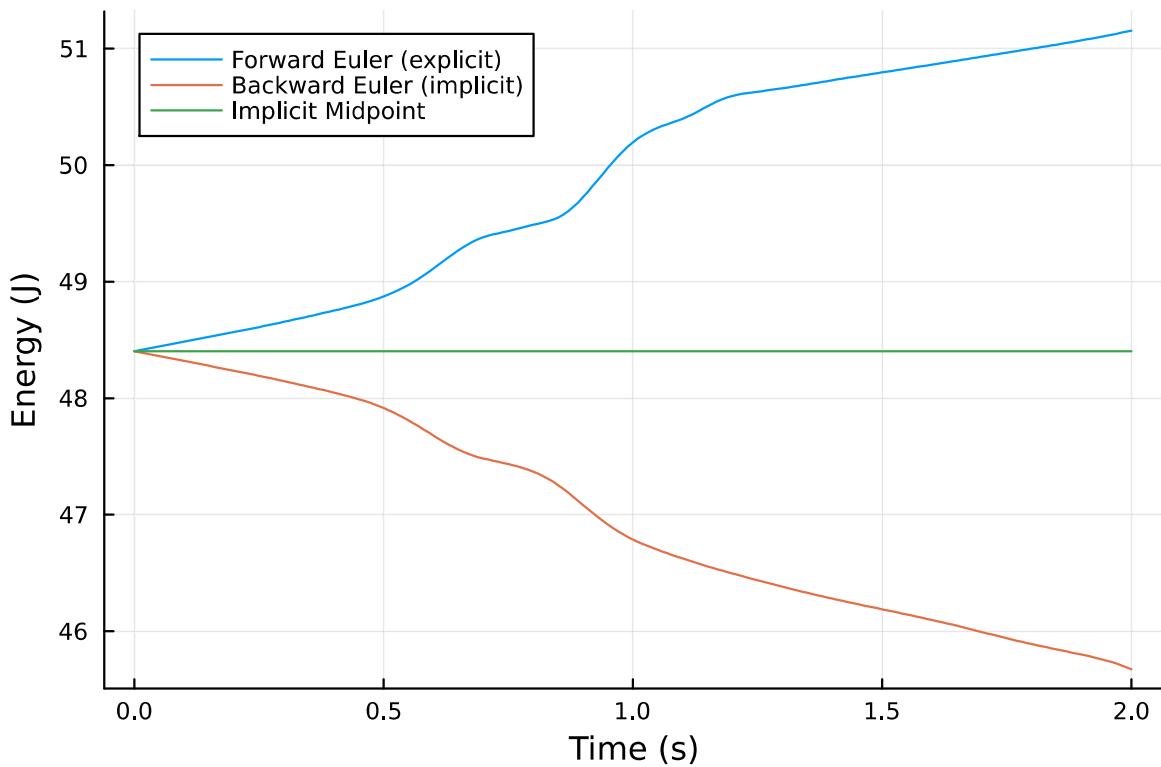
```
In [17]: @testset "energy behavior" begin
    # simulate with all integrators
    dt = 0.01
    t_vec = 0:dt:tf
    E1 = simulate_explicit(params,double_pendulum_dynamics,forward_euler,x0,dt,tf)[2]
    E2 = simulate_implicit(params,double_pendulum_dynamics,backward_euler,x0,dt,tf)
    E3 = simulate_implicit(params,double_pendulum_dynamics,implicit_midpoint,x0,dt,tf)
    E4 = simulate_implicit(params,double_pendulum_dynamics,hermite_simpson,x0,dt,tf)
    E5 = simulate_explicit(params,double_pendulum_dynamics,midpoint,x0,dt,tf)[2]
    E6 = simulate_explicit(params,double_pendulum_dynamics,rk4,x0,dt,tf)[2]

    # plot forward/backward euler and implicit midpoint
    plot(t_vec,E1, label = "Forward Euler (explicit)")
    plot!(t_vec,E2, label = "Backward Euler (implicit)")
    display(plot!(t_vec,E3, label = "Implicit Midpoint", xlabel = "Time (s)", ylabel = "Energy Error"))

    # test energy behavior
    E0 = E1[1]

    @test 2.5 < (E1[end] - E0) < 3.0
    @test -3.0 < (E2[end] - E0) < -2.5
    @test abs(E3[end] - E0) < 1e-2
    @test abs(E0 - E4[end]) < 1e-4
    @test abs(E0 - E5[end]) < 1e-1
```

```
@test abs(E0 - E6[end]) < 1e-4
end
```



**Test Summary:** | Pass Total Time  
energy behavior | 6 6 0.1s

Test.DefaultTestSet("energy behavior", Any[], 6, false, false, true, 1.738638690919e9, 1.738638691014e9, false, "c:\\Users\\zsq4\\Desktop\\OCRL\\Optimal-Control-and-Reinforcement-Learning\\HW1\_S25\\j1\_notebook\_cell\_df34fa98e69747e1a8f8a730347b8e2f\_X33sZmlsZQ==.jl")

Another important takeaway from these integrators is that explicit Euler results in unstable behavior (as shown here by the growing energy), and implicit Euler results in artificial damping (losing energy). Implicit midpoint however maintains the correct energy. Even though the solution from implicit midpoint will vary from the initial energy, it does not move secularly one way or the other.